Xinyi (Cindy) Zhou

213-272-3352 xinyi.zhou223@gmail.com

www.linkedin.com/in/xinyi-zhou-cindy/ github.com/CindyChow123 cindychow123.github.io/

Education

University of Southern California **January 2023 – December 2024**
M.S. in Computer Science; GPA: 4.0/4.0 *Los Angeles, CA*

Southern University of Science and Technology **September 2018 – July 2022**
B.E. in Computer Science and Technology; ***Outstanding Graduate Award****; GPA: 3.75/4.0* *Shenzhen, China*

Technical Skills

Languages: TypeScript, Python, SQL, JavaScript, Java, C#, C++, HTML/CSS
Backend/Frameworks: REST APIs, FastAPI, Spring Boot, XMPP; React, Vue.js
Data/Infra: BigQuery, Snowflake, PostgreSQL, Amplitude, ETL; AWS, Docker, CI/CD; Better Stack
ML/LLM: PyTorch, Hugging Face Transformers, DeepSpeed, Prompt Engineering, LLM Evaluation

Work Experience

Software Engineer (Full Stack) | *Channel AI* | *Menlo Park, US* **Feb 2025 – Present**

- Led the development and rollout of a **personalization feature**, delivering **+45% revenue uplift** for mature cohorts within 4 days (Snowflake-validated) and a **~4% increase** in daily messaging in the treatment cohort.
- Developed and maintained backend APIs for iOS/Android/web clients; optimized SQL, XMPP request flows, and LLM inference calls to internal model servers to improve latency, reliability, and scalability.
- Built an hourly **TypeScript** ETL pipeline on BigQuery (Amplitude exports) to refresh recommendation training data/features; added retries and windowed backfills to handle late-arriving events and maintain freshness.
- Built internal Trust & Safety tooling (text/image moderation + quality checks) integrating third-party classifiers; reduced manual review from 100% to **22%** by routing only **high-risk, policy-flagged** items to human review.
- Improved conversational and image-generation consistency across AI characters by iterating on prompts and building an **automated evaluation** loop (quantitative pass/fail tests + LLM-driven prompt refinement).
- Improved the company website's SEO via technical and on-page optimizations (metadata, structured content, performance, crawlability) to increase search discoverability.

Machine Learning (NLP) Engineer Intern | *UBTECH* | *Shenzhen, China* **May 2023 – August 2023**

- Fine-tuned LLM Vicuna chatbots (7B, 13B, and 33B) with **DeepSpeed ZeRO** and **LoRA** on a V100 GPU cluster for more engaging human-robot interactions in reception and tour guiding, obtaining an increase of 29.74% in user preference and 22 in Elo rating.
- Collected over 50,000 high-quality ChatGPT-generated receptioning and guiding dialogues through **Python API**, cleaned and transformed them into JSON files for storage using **Python**.
- Deployed a Chatbot implemented by **Gradio**, **Uvicorn**, and **FastAPI** on a cloud service for users to evaluate fine-tuned models.

Publications

Xinyi Zhou, Zeinabsadat Saghi, Sadra Sabouri, Rahul Pandita, Mollie McGuire, Souti Chattopadhyay. "Cognitive Biases in LLM-Assisted Software Development." *ICSE '26*, 2026. arXiv:2601.08045

Sadra Sabouri, Philipp Eibl, **Xinyi Zhou**, Morteza Ziyadi, Nenad Medvidovic, Lars Lindemann, Souti Chattopadhyay. "Trust Dynamics in AI-Assisted Development: Definitions, Factors, and Implications." *ICSE '25*, 2025. doi:10.1109/ICSE55347.2025.00199

Projects

Cognitive Biases in LLM-Assisted Software Development | *Human–AI Interaction Researcher* **January 2024 – October 2025**

- Primary researcher for an ICSE 2026 paper on cognitive biases in LLM-assisted software development; led literature review, study design, recruitment, execution, and synthesis (arXiv:2601.08045).
- Ran mixed-methods research with developers (observations: n=14; surveys: n=22), combining behavioral coding and statistical analysis to measure bias in human–LLM workflows.
- Built a validated taxonomy of 15 bias categories (systematic analysis of 90 cognitive biases) and translated findings into concrete product/engineering guidelines for designing and evaluating LLM-driven systems with bias-mitigation guardrails.
- Measured bias prevalence in human–LLM workflows (48.8% biased actions; 56.4% tied to developer–LLM interactions), using results to define guardrail requirements and evaluation metrics.

Color Correction WeChat Mini Program | *Front-end Developer, Project Manager* **February 2021 – May 2021**

- Implemented the front-end using **HTML/CSS** and developed a **JavaScript** algorithm for color space conversion that simulated the visual experience of individuals with color blindness to promote equality for them.
- Collaborated with Tencent Charity Foundation to write product requirement documents and managed team member's development schedule to ensure the product was finished on time